




Global state

Johan Montelius and Vladimir Vlassov





Global state

Time is very much related to the notion of *global state*.


If we cannot agree on a time, how should we agree on a global state?

Global state is important:

- Garbage collection
- Dead-lock detection
- Termination
- Debugging

Global state


Given a partial order of events, can we say anything about the state of the system?





Local history and Local state

The **history** of a process is a sequence of events: $\langle p_0, p_1, \dots, p_n \rangle$



The **state** of a process is a description of the process after (before) an event.

- A state corresponds to a **finite prefix of the process's history**



Global history and Global state

What is the **global history** of concurrent distributed processes?

- The union of individual histories of all processes?
- Do all unions make sense?

What is the **global state** of a distributed system?


- The union of states of individual processes?

A **global state** corresponds to initial prefixes of the individual process histories.




Global history and Cut

A **cut** is a subset in the global history up to a specific event in each history.




An event is in the *cut* if it belongs to the events of a history up to the specific event.

All cuts are equal, but ...






..some are more equal ..





.. than others






Consistent cuts

For each event e in the cut:

- if f happened before e then
- f is also in the cut.






Consistent global state

A **consistent cut** corresponds to a **consistent global state**.

- it is a **possible state** without contradictions
- it is **consistent with** the **actual execution**
- the actual execution might not have passed through the state, even though it's consistent

Consistent, but not actual states




All real time cuts are consistent, but who knows the real time?



Linearization


- A **run** is a total ordering of all events in a global history that is consistent with each local history.
- A **linearization** or **consistent run** is a run that describes transitions between **consistent global states**.
- A state S' is **reachable** from state S if there is a linearization from S to S' .

Linearization




Possible state transitions

[p0, p1, q0, r0, q1, r1, p2, p3, q2, r2, q3]




Possible state transitions

[q0, p0, p1, r0, q1, r1, p2, p3, q2, r2, q3]



Possible paths



Each path is a consistent run, a linearization, one of which the execution actually took.



Why is this important?

- If we can collect all events and know the happened before order, then we can construct all possible linearizations.
- We know that the actual execution took one of these paths.
- Can we say something about the execution even though we do not know which path that was taken?



Global state predicate

A global state predicate is a property that is true or false for a global state.


- **Safety** - a predicate is never (or always) true in any state.
- **Liveness** - a predicate that eventually evaluates to true.

How do we determine if a property holds in an execution?




Let's capture all linearizations

Idea - use **vector clocks**, collect all events of the execution.



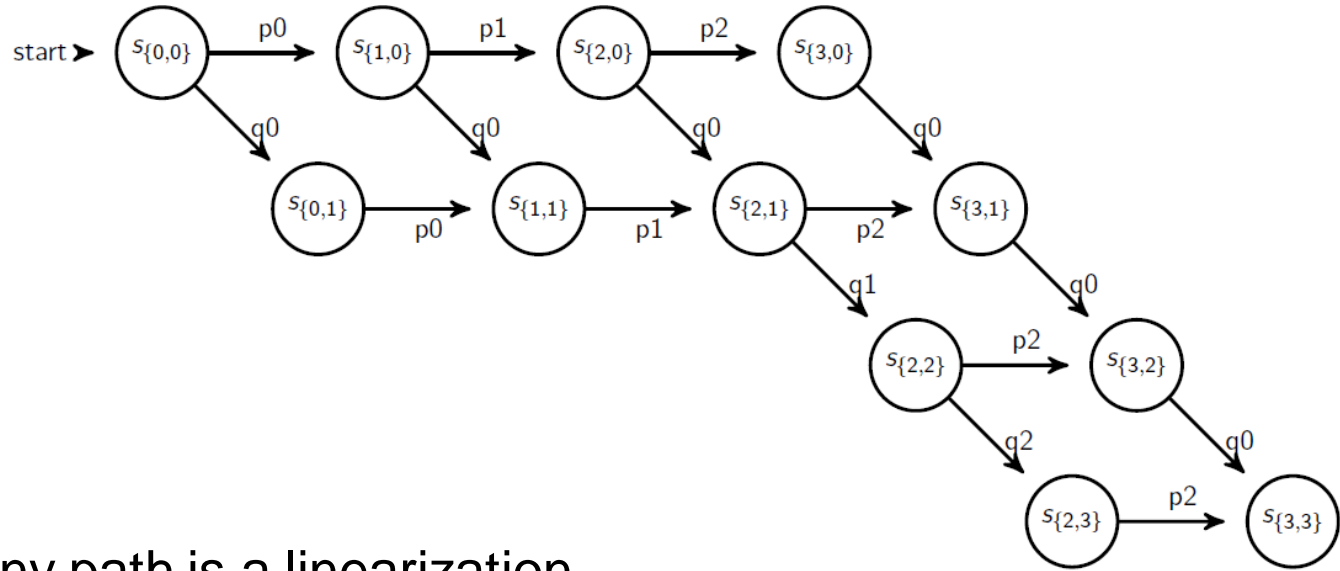


Construct all linearizations





An execution lattice

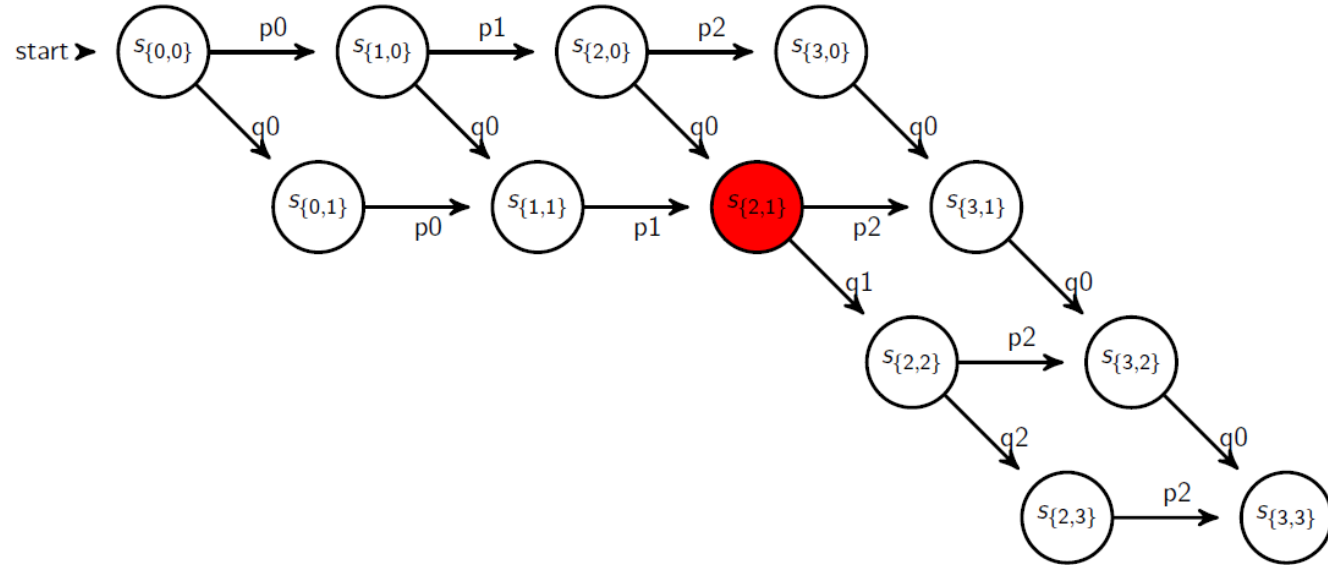


Any path is a linearization.

The actual execution took one path.

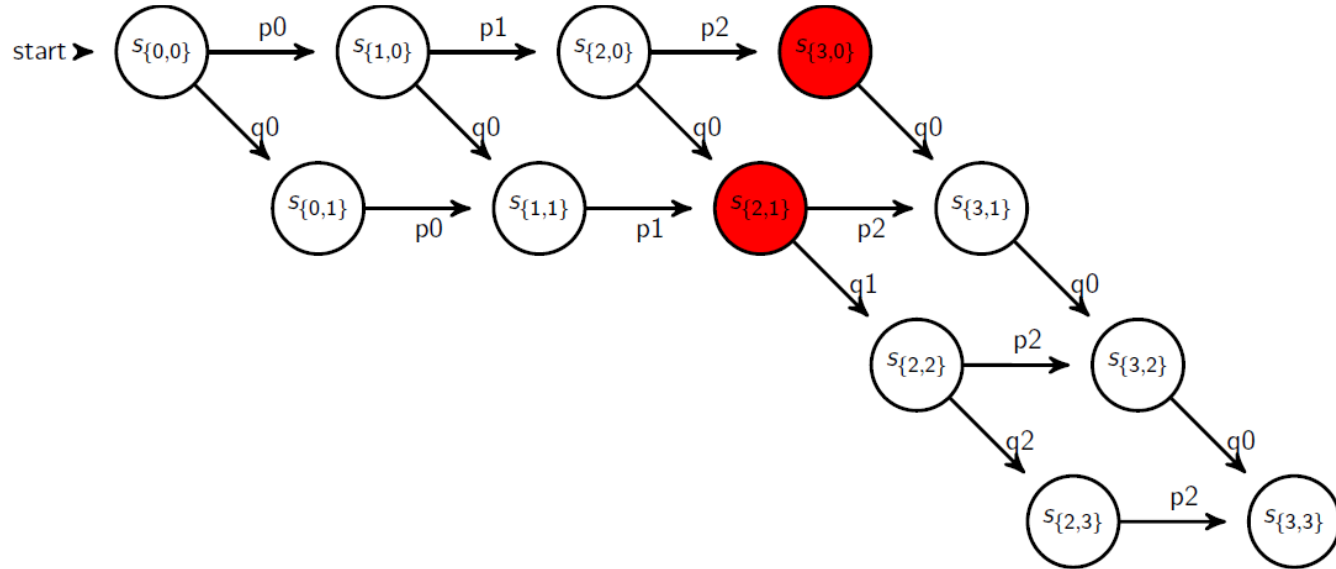


Possibly true



If a predicate is true in a consistent global state of the lattice, then it is *possibly true* in the execution.

Definitely true



If we cannot find a path from the initial state to the final state without reaching a state for which a predicate is true then the predicate is *definitely true* during the execution.

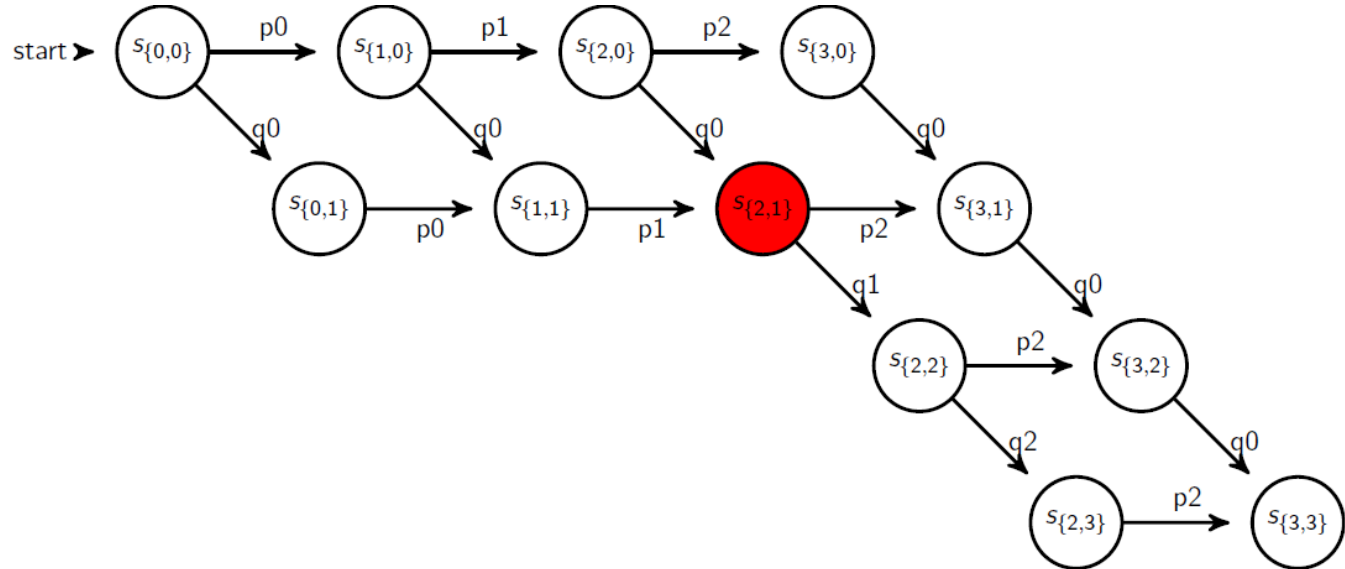
Stable and non-stable

We differentiate between:

- **Stable**: if a predicate is true it remains true for all reachable states
- **Non-stable**: if a predicate can become true and then later become false

Stable is good

What do I know if a stable predicate is true for state $S_{\{2,1\}}$?





Let's capture a possible state

Idea: capture a consistent global state that was possibly true in the execution.

If a stable predicate is true for this state - then it is true in the actual execution.

How do we capture a state?


Snapshot - Chandy and Lamport

A node initiates a snapshot when it receives a *marker*.

- Record the local state and
- send a *marker* on all out going channels.
- Record all incoming messages on each channel, ..
- until you receive a marker.
- When the last channel is closed you have a local and a set of messages.

Ask one node to initiate the snapshot, collect all local states and messages and construct a global state.


Snapshot markers



What messages are collected by which node?



Snapshot markers



What messages are collected by which node?



Snapshot

- Allows us to collect a global state during execution.
- Only allows us to determine stable predicates.



Summary

The happened before order gives us *consistent cuts or consistent global states*.

Using vector clocks we can time stamp states, *construct all possible linearizations* and evaluate if predicates hold true in the execution.

A snapshot can record a consistent state that can be used to evaluate *stable predicates*.